

VisionFood QAI — AI-Powered Quality Inspection for Food & Beverage Manufacturing

Project Review 2: Implementation & Architecture

Shreyas Bairy K S – 1BM23CD058

Sushanth S – 1BM23CD063

Ullas N – 1BM23CD067

B.M.S. College of Engineering, Bengaluru-560019
(Autonomous Institute, Affiliated to VTU)

Department of Computer Science and Engineering (Data Science)

Guide: Prof. Nayana N Kumar

Project Phase 1 [AY: 2025–26]

Outline

- 1 Literature Survey
- 2 Technical Details: Uncertainty Quantification and Remediation
- 3 References

Literature Survey — Research Domains

We reviewed **16 peer-reviewed papers** (2024–2025) spanning six research domains in deep learning and industrial quality inspection:

- ① CNN-based Defect Detection
- ② YOLO Object Detection
- ③ Food Quality Assessment
- ④ Industrial Vision Systems
- ⑤ IoT & Sensor Fusion
- ⑥ Transfer Learning & Anomaly Detection
- ⑦ Systematic Review & Framework

Key insight: The literature converges on multi-stage pipelines combining detection + classification, with emerging focus on **explainability**, **incremental learning**, and **supply-chain integration**.

Sources: IEEE ICCAS, ICCSP, GCAT; Elsevier; Vietnam Journal of Computer Science; Foods; J. Imaging

Domain 1 & 2: CNN-Based Defect Detection & YOLO Object Detection

Key Finding: Real-Time Multi-Defect Detection is Feasible

Papers by Chatchapol et al. (ICCAS 2024), Kavitha & Mamatha (IJISRT 2024, GCAT 2024), and Nguyen et al. (VJCS 2025) demonstrate:

- **YOLOv5/v8 achieves >90% accuracy** on real production lines for bottle caps, liquid level, and packaging defects
- **Single-pass detection** (one-stage networks) is preferred over two-stage methods for speed
- **Hardware integration proven:** IoT sensors + PLCs (S7-1200) + embedded systems (Jetson NX) enable end-to-end automation

Highlighted papers:

- Chatchapol et al. — CNN for liquid fill (98.89% accuracy) + conveyor speed control
- Nguyen & Mac — YOLOv8 + multi-threading for real-time classification; scratch defects hardest to detect
- Kavitha et al. — YOLOv5 + Arduino actuation for multi-material sorting (cloth, PVC, PCB)

Domain 3 & 4: Food Quality Assessment & Industrial Vision Systems

Key Finding: Domain-Specific and Comprehensive Multi-Parameter Inspection

Farhangi et al. (CCDS 2024), Sanjay (ICCSP 2024), and Du (MMIA 2025) show:

- **Multi-parameter inspection** (fill level + cap + label + surface) achievable in single pass: **95.6% average accuracy**
- **Food-specific defects** (oil quality, contamination) require specialized algorithms; classical vision + deep learning are complementary
- **Online (real-time) inspection outperforms offline:** 98% accuracy vs. 85–92%, 85% faster processing (3 sec vs. 20 sec per unit)

Highlighted papers:

- Farhangi et al. — LabVIEW + classical vision; 100% cap accuracy, 95% label placement
- Sanjay — CNN for food oil quality; real-time cost-effective alternative to lab testing
- Du — Multi-sensor fusion + edge computing; 60% efficiency gain, 4.85% scrap reduction

Domains 3-7: Food Quality, Industrial Vision, IoT, Transfer Learning, & Systematic Review

Key Findings Across Remaining Domains

- **Food Quality Assessment:** Multi-parameter inspection (fill + cap + label + surface) achievable in single pass with 95.6% accuracy (Farhangi et al.); specialized algorithms for food-specific defects (oil, contamination) (Sanjay)
- **Industrial Vision:** Online real-time inspection outperforms offline: 98% accuracy vs. 85–92%, 85% faster processing (Du); incremental learning enables adaptation to new defect types without full retraining (Ribeiro et al.)
- **IoT & Sensor Fusion:** RFID + LSTM enables item-level spoilage prediction (95.2% accuracy); supply-chain end-to-end monitoring (Cao et al.; Krishnan et al.)
- **Transfer Learning & Anomaly Detection:** Autoencoder-based approaches require only clean samples; knowledge transfer across product types with 3–5 images (Truong & Luong)
- **Systematic Review (Liakos et al., PRISMA 2025):** Fresh produce dominates (36%); beverages only 4% → significant research opportunity. Emerging trends: hyperspectral imaging, explainable AI, blockchain traceability.

Primary Base Paper — Liu et al. (2025)

Research on Mechanical Automatic Food Packaging Defect Detection Based on Improved YOLOv5

Liu et al., PLOS ONE (2025)

Precision: 0.96 | Recall: 0.94 | F1-Score: 0.94

Dataset: 7,400 food packaging images, 5 defect classes

Why Liu et al. is our foundation:

- **Direct relevance:** Multi-defect detection on food packaging using YOLO-based deep learning
- **Proven methodology:** Achieves 94% F1-score on real production data
- **Defect coverage:** 5 classes (scratches, dents, stains, deformations, label errors) — 4 map directly to our taxonomy

Our project extends this work with:

- YOLOv8-based detector for stronger real-time defect detection
- MC Dropout uncertainty quantification for explainable, human-interpretable verdicts


Project Objectives

- Design and deploy a real-time AI-powered quality inspection system for food and beverage manufacturing
- Detect and classify four critical defect types: improper filling, packaging damage, label misalignment, and surface contamination
- Achieve inference latency suitable for production lines (edge-compatible; target: <500 ms per product)
- Integrate with production line hardware (cameras, sensors, PLCs) for automated decision-making

Project Objectives (continued)

- Provide human-interpretable explanations for all automated defect decisions
- Implement uncertainty quantification to escalate low-confidence cases to human review
- Enable incremental learning so the system adapts to new defect patterns without full retraining

Implementation — Food and Beverage Defect Detection Pipeline




Implementation — Flowchart Terms (Part 1)

Terms used in the first flowchart

- **User classifies product:** the operator selects whether the item is a food product or a beverage.
- **Product type:** the high-level branch that activates the correct pipeline.
- **Detect product pack / Detect bottle:** the first detection stage that finds the product in the image.
- **Crop product pack / Crop bottle top:** isolates the region that needs finer inspection.
- **Parallel inspection:** tear detection and label detection run at the same time.
- **Run tear detector:** checks food packages for packaging damage, mainly tears.
- **Run cap detector:** checks beverage bottles for cap presence.
- **Run cap classifier:** classifies the detected cap as *good_cap*, *defective_cap*, or *no_cap*.
- **Run label detector:** checks whether the label is present and positioned correctly.
- **Label evaluator:** reserved for the user-provided label-matching logic; left blank for now.

Implementation — Output Combination and Evaluation Pipeline



Implementation — Flowchart Terms (Part 2)

Terms used in the second flowchart

- **Combined defect evidence:** the outputs from food or beverage detection are merged into one evaluation stream.
- **MC Dropout uncertainty quantification:** estimates model uncertainty before making a final judgment.
- **Explainability payload:** the evidence shown to the user, including boxes, scores, and reasons.
- **Severity Scorer:** assigns the defect a severity grade based on risk, confidence, area, and retries.
- **Severity grade:** the REMEDY class assigned to the defect, such as S1, S2, S3, or S4.
- **Triage Router:** maps the severity grade to the corrective action.
- **Remediation action:** RELABEL, REFILL, REPACK, or CLEAN for manageable defects.
- **Reject / quarantine:** used for severe or unsafe defects that should not continue on the line.
- **Persist inspection result:** saves the final decision for traceability.
- **Live WebSocket event:** sends the inspection result to the dashboard in real time.
- **Operator override:** allows a human to revise the verdict if needed.

Uncertainty Quantification — MC Dropout Theory

Monte Carlo Dropout estimates model **epistemic uncertainty** (model knowledge gaps) at inference time by keeping dropout layers active:

MC Dropout Methodology

- ➊ **Multiple forward passes:** run the classifier n times on the same input (default: $n = 20$)
- ➋ **Stochastic variation:** each pass applies different random dropout masks
- ➌ **Aggregate outputs:** collect all n predicted confidences for the defect class
- ➍ **Compute statistics:**
 - Mean confidence: $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n c_i$
 - Standard deviation (uncertainty): $\hat{\sigma} = \sqrt{\frac{1}{n} \sum_{i=1}^n (c_i - \hat{\mu})^2}$
 - 95% confidence interval: $[\hat{\mu} - 1.96\hat{\sigma}, \hat{\mu} + 1.96\hat{\sigma}]$
- ➎ **Flag uncertain predictions:** if $\hat{\sigma} > \text{threshold (0.15)}$, mark as *uncertain*

Integration with EfficientViT Classifier

- **Model export:** classifier exported to ONNX format *with dropout enabled* (using `training=True` mode)
- **Inference flow:**
 - 1 Extract product crop from detection bounding box
 - 2 Run ONNX session 20 times on the same crop image
 - 3 Collect all 20 confidence scores $\{c_1, c_2, \dots, c_{20}\}$
 - 4 Compute mean, std, and confidence interval from samples
- **Uncertainty thresholds** (tunable via `EdgeConfig`):
 - UQ threshold = 0.15: if $\hat{\sigma} > 0.15$, mark *is_uncertain* = True
 - Escalation threshold = 0.60: if mean confidence < 0.60 *AND* uncertain, escalate to fog tier
- **Result stored in** `UQResult` schema with full statistics for explainability

REMEDY Engine — Architecture Overview

REMEDY (Rapid Evaluation, Mitigation, and Diagnosis of Yield Errors) routes defects to corrective actions:

Input	Component	Output
Detection + UQ	<i>SeverityScorer</i>	Severity Grade (S1–S4)
Grade + Class	<i>TriageRouter</i>	Remediation Action
SKU ID	<i>SKUProfileManager</i>	Config Overrides

Severity Grades

- **S1 (Minor)**: remediable on-line at dedicated station (e.g., relabel, clean)
- **S2 (Moderate)**: requires intermediate handling (e.g., refill, repack)
- **S3 (Severe)**: rejected; does not re-enter production
- **S4 (Critical)**: quarantined; flagged for investigation

REMEDY Severity Scorer — Weighted Formula

Severity score (range [0, 1]) is computed as a weighted sum of four risk components:

$$\text{score} = w_{\text{area}} \cdot a + w_{\text{conf_uq}} \cdot u + w_{\text{class}} \cdot r + w_{\text{attempt}} \cdot p$$

Component Definitions

- a = **Area ratio**: defect bounding box area / total image area (capped at 0.25)
- u = **Uncertainty component**: MC Dropout std dev $\hat{\sigma}$ (capped at 0.30)
- r = **Class risk**: food safety impact of defect class (e.g., contamination=1.0, misalignment=0.30)
- p = **Attempt penalty**: prior remediation failures (capped at 1.0, increases by 0.5 per attempt)

Weights (from EdgeConfig)

$$w_{\text{area}} = 0.35, \quad w_{\text{conf_uq}} = 0.15, \quad w_{\text{class}} = 0.40, \quad w_{\text{attempt}} = 0.10$$

Triage Routing — Defect-to-Action Mapping

Routing Table: (Defect Class, Severity Grade) → Remediation Action

Defect Class	Grade	Action	Station
Label Misalignment	S1/S2	RELABEL	Station A
Improper Filling	S1/S2	REFILL	Station B
Packaging Damage	S1/S2	REPACK	Station C
Surface Contamination	S1	CLEAN	Station C
Any class	S3/S4	REJECT	Quarantine
Max attempts exceeded	Any	REJECT	Quarantine
No defects	—	PASS	Shipping

Action Definitions and Station Assignments

- **RELABEL**: re-apply or correct label (Station A)
- **REFILL**: adjust product level or contents (Station B)
- **REPACK**: replace damaged packaging (Station C)
- **CLEAN**: surface decontamination (Station C)
- **REJECT**: remove from line permanently (Quarantine)
- **PASS**: product proceeds to shipping (no intervention)

Each remediation action is routed to its corresponding station, where operators perform the corrective measure. Severe defects (S3/S4) and products exceeding max remediation attempts are automatically rejected, bypassing all stations.

Product Customization — SKU Profiles

SKU Profile Manager loads per-product YAML configurations to override REMEDY parameters:

YAML Configuration Structure

- **class_risk_overrides**: per-class risk weights (e.g., contamination: 0.95, filling: 0.80)
- **rejection_area_thresholds**: maximum defect size before forced rejection per class
- **preferred_stations**: hint for routing (e.g., bottle_250ml prefers Station B for filling checks)
- **uq_threshold_override**: product-specific uncertainty escalation threshold

Example: bottle_250ml.yaml

```
sku_id: "bottle_250ml"
class_risk_overrides:
  improper_filling: 0.85 # Higher risk
  label_misalignment: 0.35
rejection_area_thresholds:
  packaging_damage: 0.10
preferred_stations:
  improper_filling: "B"
```

References



B. Chatchapol, W. Autanan, and W. Anan, "Liquid filling quality control system for beverage industry using image processing and cnn," in *2024 24th International Conference on Control, Automation and Systems (ICCAS)*, pp. 1–6, 2024.



K. K. S and M. C. G, "Automated product defect detection using image processing techniques for effective sorting and quality assurance: A survey," *International Journal of Innovative Science and Research Technology (IJISRT)*, vol. 9, no. 6, 2024.



S. Sanjay et al., "Food oil quality and usage detection using ai," in *2024 10th International Conference on Communication and Signal Processing (ICCSP)*, pp. 1190–1194, 2024.



K. K. S, M. C. G, and R. T. V, "Deep learning based product defect detection for sustainable smart manufacturing," in *2024 5th IEEE Global Conference for Advancement in Technology (GCAT)*, pp. 1–6, 2024.



O. Farhangi, E. Sheidaee, and A. Kisalaei, "Machine vision for detecting defects in liquid bottles: An industrial application for food and packaging sector," *Cluster Computing and Data Science*, 2024.



X.-T. Nguyen, T.-T. Mac, Q.-D. Nguyen, and H.-A. Bui, "An industrial system for inspecting product quality based on machine vision and deep learning," *Vietnam Journal of Computer Science*, vol. 12, no. 2, pp. 193–208, 2025.



P. Krishnan, U. S. Ebenezer, R. Ranitha, N. Purushotham, and T. S. Balakrishnan, "Ai-driven intelligent iot systems for real-time food quality monitoring and analysis," in *2024 International Conference on Trends in Quantum Computing and Emerging Business Technologies*, pp. 1–5, 2024.



A. M. Truong and H. Q. Luong, "A non-destructive, autoencoder-based approach to detecting defects and contamination in reusable food packaging," *Current Research in Food Science*, vol. 8, p. 100758, 2024.



T. Du, "Online measurement and detection technology and its optimized application in intelligent manufacturing environment," *Procedia Computer Science*, vol. 262, pp. 83–91, 2025.



Z. Cao, Z. Wu, and J. Gray, "Rfid tag-integrated multi-sensors with aiot cloud platform for food quality analysis," *Electronics*, vol. 15, no. 1, 2025.

References (cont.)



K. G. Liakos, V. Athanasiadis, E. Bozinou, and S. I. Lalas, "Machine learning for quality control in the food industry: A review," *Foods*, vol. 14, 2025.



A. G. Ribeiro, L. Vilaça, C. Costa, T. S. da Costa, and P. M. Carvalho, "Automatic visual inspection for industrial application," *Journal of Imaging*, vol. 11, 2025.



R. R. Subramanian, M. Dhamini, M. Srija, M. Vaishnavi, and M. Vidyadhari, "Revolutionizing manufacturing quality control with optimized resnet50: A deep learning approach for real-time defect detection," in *2025 International Conference on Computational Robotics, Testing and Engineering Evaluation (ICCRTEE)*, 2025.



F. Xiong, N. Kühn, and M. Stauder, "Designing a computer-vision-based artifact for automated quality control: a case study in the food industry," *Annals of Operations Research*, 2024.



N. Banús, I. Boada, P. Xiberta, P. Toldrà, and N. Bustins, "Deep learning for the quality control of thermoforming food packages," *Scientific Reports*, vol. 11, 2021.



N. C. Gediya, "Smart manufacturing: Real-time quality control with ai and image recognition," *Conference Proceedings*, 2024.



H. Lawrence, B. Pearce, P. Resnick, *et al.*, "Implicit bias of linear equivariant networks," in *Proceedings of the 39th International Conference on Machine Learning (ICML)*, vol. 162 of *Proceedings of Machine Learning Research*, pp. 12096–12125, 2022.



G. Liu, S. Zhang, L. Wang, X. Li, and G. Li, "Research on mechanical automatic food packaging defect detection model based on improved yolov5 algorithm," *PLOS ONE*, vol. 20, 2025.

Thank You

Questions & Discussion